

```
In [3]: import tensorflow as tf
        from tensorflow.keras import datasets, layers, models
        import numpy as np
        import matplotlib.pyplot as plt

In [4]: (X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()

In [5]: X_test.shape

Out[5]: (10000, 32, 32, 3)

In [6]: X_train.shape

Out[6]: (50000, 32, 32, 3)

In [7]: y_test.shape

Out[7]: (10000, 1)

In [8]: y_train.shape

Out[8]: (50000, 1)

In [9]: y_train[:5]

Out[9]: array([[6],
               [9],
               [9],
               [4],
               [1]], dtype=uint8)

In [10]: y_train = y_train.reshape(-1,)
         y_train[:5]


Out[10]: array([6, 9, 9, 4, 1], dtype=uint8)

In [11]: y_test = y_test.reshape(-1,)

In [12]: classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']


In [13]: def plot_sample(X, y, index):
         plt.figure(figsize=(15,2))
         plt.imshow(X[index])
         plt.xlabel(classes[y[index]])

In [14]: plot_sample(X_train, y_train, 5)
```



```
In [15]: plot_sample(X_train, y_train, 5)

In [16]: plot_sample(X_train, y_train, 501)
```



```
In [17]: X_train = X_train / 255.0
         X_test = X_test / 255.0

In [18]: ann = models.Sequential([
         layers.Flatten(input_shape=(32,32,3)),
         layers.Dense(3000, activation='relu'),
         layers.Dense(1000, activation='relu'),
         layers.Dense(10, activation='softmax'),
         ])

         ann.compile(optimizer='SGD',
                     loss='sparse_categorical_crossentropy',
                     metrics=['accuracy'])

         ann.fit(X_train, y_train, epochs=5)

E:\Anaconda ML\lib\site-packages\keras\src\layers\reshaping\flatten.py:37: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model instead.
  super().__init__(**kwargs)

Epoch 1/5 ----- 55s 33ms/step - accuracy: 0.3040 - loss: 1.9352
Epoch 2/5 ----- 55s 35ms/step - accuracy: 0.4225 - loss: 1.6385
Epoch 3/5 ----- 51s 33ms/step - accuracy: 0.4531 - loss: 1.5545
Epoch 4/5 ----- 50s 32ms/step - accuracy: 0.4737 - loss: 1.4950
Epoch 5/5 ----- 48s 30ms/step - accuracy: 0.4953 - loss: 1.4420

Out[18]: <keras.src.callbacks.history.History at 0x1e6ed7663d0>
```

```
In [19]: from sklearn.metrics import confusion_matrix, classification_report
         import numpy as np
         y_pred = ann.predict(X_test)
         y_pred_classes = [np.argmax(element) for element in y_pred]

         print('classification report: \n', classification_report(y_test, y_pred_classes))

313/313 ----- 2s 7ms/step
classification report:
              precision    recall  f1-score   support


     0         0.52        0.60         0.56       1000
     1         0.66        0.53         0.58       1000
     2         0.49        0.12         0.19       1000
     3         0.38        0.19         0.25       1000
     4         0.39        0.45         0.42       1000
     5         0.43        0.32         0.37       1000
     6         0.59        0.38         0.46       1000
     7         0.32        0.79         0.45       1000
     8         0.58        0.64         0.61       1000
     9         0.52        0.60         0.56       1000

 accuracy          0.49          0.46          0.46      10000
 macro avg         0.49          0.46          0.45      10000
weighted avg         0.49          0.46          0.45      10000
```

```
In [20]: import seaborn as sns

In [21]: plt.figure(figsize=(14,7))
         sns.heatmap(y_pred, annot = True)
         plt.ylabel('Truth')
         plt.xlabel('Prediction')
         plt.title('Confusion matrix')
         plt.show

Out[21]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [22]: cnn = models.Sequential([
         layers.Conv2D(filters=32, kernel_size=(3,3), activation='relu', input_shape=(32,32, 3)),
         layers.MaxPooling2D((2,2)),

         layers.Conv2D(filters=64, kernel_size=(3,3), activation='relu'),
         layers.MaxPooling2D((2,2)),

         layers.Flatten(),
         layers.Dense(64, activation='relu'),
         layers.Dense(10, activation='softmax'),
         ])

E:\Anaconda ML\lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

In [23]: cnn.compile(optimizer='adam',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

In [24]: cnn.fit(X_train, y_train, epochs=10)

Epoch 1/10 ----- 13s 7ms/step - accuracy: 0.3983 - loss: 1.6582
Epoch 2/10 ----- 11s 7ms/step - accuracy: 0.6119 - loss: 1.1124
Epoch 3/10 ----- 11s 7ms/step - accuracy: 0.6695 - loss: 0.9481
Epoch 4/10 ----- 11s 7ms/step - accuracy: 0.7075 - loss: 0.8428
Epoch 5/10 ----- 11s 7ms/step - accuracy: 0.7288 - loss: 0.7750
Epoch 6/10 ----- 11s 7ms/step - accuracy: 0.7536 - loss: 0.7110
Epoch 7/10 ----- 12s 7ms/step - accuracy: 0.7706 - loss: 0.6543
Epoch 8/10 ----- 12s 8ms/step - accuracy: 0.7877 - loss: 0.6031
Epoch 9/10 ----- 11s 7ms/step - accuracy: 0.8054 - loss: 0.5548
Epoch 10/10 ----- 12s 8ms/step - accuracy: 0.8267 - loss: 0.5020

Out[24]: <keras.src.callbacks.history.History at 0x1e7049849d0>
```

```
In [25]: cnn.evaluate(X_test, y_test)

313/313 ----- 1s 3ms/step - accuracy: 0.7152 - loss: 0.9070

Out[25]: [0.9251238107681274, 0.7091000080108643]

In [26]: y_pred = cnn.predict(X_test)
         y_pred[:5]

313/313 ----- 1s 3ms/step

Out[26]: array([[3.13081604e-04, 4.44748948e-05, 5.50113909e-04, 7.99367249e-01,
               8.86134439e-05, 4.03291918e-02, 1.07765093e-01, 6.09774315e-06,
               5.11943251e-02, 3.41743929e-04],
               [5.92878573e-02, 1.65216730e-03, 1.35267464e-05, 9.33361566e-09,
               8.25674178e-08, 2.40147768e-09, 1.99592201e-08, 1.20279198e-09,
               9.38989563e-01, 1.47697487e-04],
               [5.96161075e-02, 1.73812702e-01, 2.15590112e-02, 7.67891994e-03,
               1.13394135e-03, 3.51405656e-03, 2.92896270e-03, 2.28178990e-03,
               7.09434980e-01, 1.80395674e-02],
               [9.83872950e-01, 4.68066632e-04, 6.39343983e-04, 1.09955153e-05,
               4.20502511e-05, 7.97551638e-06, 4.19331955e-05, 9.58110468e-06,
               1.47862770e-02, 1.12778129e-04],
               [7.27530676e-07, 4.49761183e-06, 7.50960410e-03, 1.68715734e-02,
               4.27279890e-01, 6.18672639e-04, 5.47628284e-01, 1.12274364e-07,
               8.65199294e-05, 7.82818829e-08]])
```

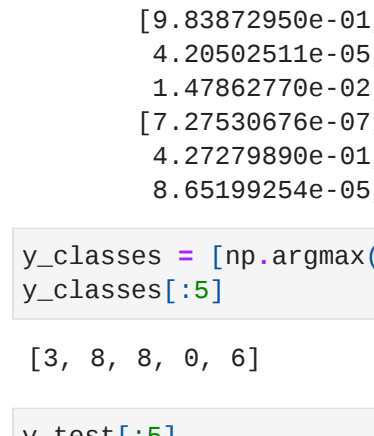
```
In [27]: y_classes = [np.argmax(element) for element in y_pred]
         y_classes[:5]

Out[27]: [3, 8, 8, 0, 6]

In [28]: y_test[:5]

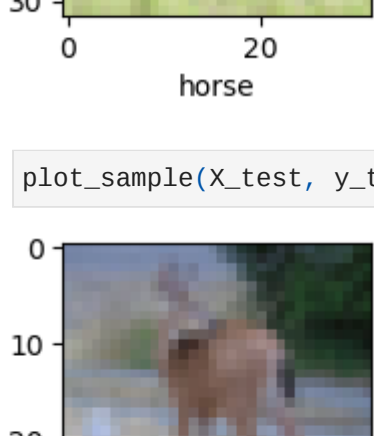
Out[28]: array([3, 8, 8, 0, 6], dtype=uint8)

In [29]: plot_sample(X_test, y_test, 60)
```



```
In [30]: plot_sample(X_test, y_test, 100)

In [37]: plot_sample(X_test, y_test, 79)
```



```
In [38]: classes[y_classes[79]]

Out[38]: 'ship'

In [ ]:
```

