

```
In [3]:import tensorflow as tf
        from tensorflow.keras import datasets, layers, models
        import numpy as np
        import matplotlib.pyplot as plt

In [4]:(X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()

In [5]:X_test.shape

Out[5]:(10000, 32, 32, 3)

In [6]:X_train.shape

Out[6]:(50000, 32, 32, 3)

In [7]:y_test.shape

Out[7]:(10000, 1)

In [8]:y_train.shape

Out[8]:(50000, 1)

In [9]:y_train[5]

Out[9]: array([[6,
[9],
[9],
[4],
[1]], dtype=uint8)

In [10]:y_train = y_train.reshape(-1)
        y_train[5]


Out[10]:array([6, 9, 9, 4, 1], dtype=uint8)

In [11]:y_test = y_test.reshape(-1)

In [12]:classes = ['airplane','automobile','bird','cat','deer','dog','frog','horse','ship','truck']


In [13]:def plot_sample(X, y, index):
        plt.figure(figsize=(15,2))
        plt.imshow(X[index])
        plt.xlabel(classes[y[index]])

In [14]:plot_sample(X_train, y_train, 5)
```



```
In [15]:plot_sample(X_train, y_train, 5)

In [16]:plot_sample(X_train, y_train, 501)
```



```
In [17]:X_train=X_train / 255.0
        X_test = X_test / 255.0

In [18]:ann = models.Sequential([
        layers.Flatten(input_shape=(32,32,3)),
        layers.Dense(3000, activation = 'relu'),
        layers.Dense(1000, activation = 'relu'),
        layers.Dense(10, activation = 'softmax'),
    ])

    ann.compile(optimizer='SGD',
                loss='sparse_categorical_crossentropy',
                metrics=['accuracy'])

    ann.fit(X_train, y_train, epochs=5)

E:\Anaconda ML\lib\site-packages\keras\src\layers\reshaping\flatten.py:37: UserWarning: Do not pass an `input_shape` / `input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
Epoch 1/5
1563/1563 — 55s 33ms/step - accuracy: 0.3040 - loss: 1.9352
Epoch 2/5
1563/1563 — 55s 35ms/step - accuracy: 0.4225 - loss: 1.6385
Epoch 3/5
1563/1563 — 51s 33ms/step - accuracy: 0.4531 - loss: 1.5545
Epoch 4/5
1563/1563 — 50s 32ms/step - accuracy: 0.4737 - loss: 1.4950
Epoch 5/5
1563/1563 — 48s 30ms/step - accuracy: 0.4953 - loss: 1.4420

Out[18]:<keras.src.callbacks.history.History at 0x1e6ed7663d0>
```

```
In [19]:from sklearn.metrics import confusion_matrix, classification_report
        import numpy as np
        y_pred = ann.predict(X_test)
        y_pred_classes = [np.argmax(element) for element in y_pred]

        print('classification report:\n', classification_report(y_test, y_pred_classes))

313/313 — 2s 7ms/step
classification report:
              precision    recall  f1-score   support


     0           0.52         0.60         0.56         1000
     1           0.66         0.53         0.58         1000
     2           0.49         0.12         0.19         1000
     3           0.38         0.19         0.25         1000
     4           0.39         0.45         0.42         1000
     5           0.43         0.32         0.37         1000
     6           0.59         0.38         0.46         1000
     7           0.32         0.79         0.45         1000
     8           0.58         0.64         0.61         1000
     9           0.52         0.60         0.56         1000

 accuracy          0.49         0.46         0.46        10000
 macro avg         0.49         0.46         0.45        10000
 weighted avg      0.49         0.46         0.45        10000
```

```
In [20]:import seaborn as sns

In [21]:plt.figure(figsize=(14,7))
        sns.heatmap(y_pred, annot = True)
        plt.xlabel('Truth')
        plt.ylabel('Prediction')
        plt.title('Confusion matrix')
        plt.show

Out[21]:<function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [22]:cnn = models.Sequential([
        layers.Conv2D(filters=32, kernel_size=(3,3), activation='relu', input_shape=(32,32, 3)),
        layers.MaxPooling2D((2,2)),

        layers.Conv2D(filters=64, kernel_size=(3,3), activation='relu'),
        layers.MaxPooling2D((2,2)),

        layers.Flatten(),
        layers.Dense(64, activation='relu'),
        layers.Dense(10, activation='softmax'),
    ])

    E:\Anaconda ML\lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape` / `input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
      super().__init__(activity_regularizer=regularizer, **kwargs)

In [23]:cnn.compile(optimizer='adam',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

In [24]:cnn.fit(X_train, y_train, epochs=10)

Epoch 1/10
1563/1563 — 13s 7ms/step - accuracy: 0.3983 - loss: 1.6582
Epoch 2/10
1563/1563 — 11s 7ms/step - accuracy: 0.6119 - loss: 1.1124
Epoch 3/10
1563/1563 — 11s 7ms/step - accuracy: 0.6695 - loss: 0.9481
Epoch 4/10
1563/1563 — 11s 7ms/step - accuracy: 0.7075 - loss: 0.8428
Epoch 5/10
1563/1563 — 11s 7ms/step - accuracy: 0.7288 - loss: 0.7750
Epoch 6/10
1563/1563 — 11s 7ms/step - accuracy: 0.7536 - loss: 0.7110
Epoch 7/10
1563/1563 — 12s 7ms/step - accuracy: 0.7706 - loss: 0.6543
Epoch 8/10
1563/1563 — 12s 8ms/step - accuracy: 0.7877 - loss: 0.6031
Epoch 9/10
1563/1563 — 11s 7ms/step - accuracy: 0.8054 - loss: 0.5548
Epoch 10/10
1563/1563 — 12s 8ms/step - accuracy: 0.8267 - loss: 0.5020

Out[24]:<keras.src.callbacks.history.History at 0x1e7049849d0>
```

```
In [25]:cnn.evaluate(X_test, y_test)

313/313 — 1s 3ms/step - accuracy: 0.7152 - loss: 0.9070

Out[25]:[0.9251238107681274, 0.7091000080108643]

In [26]:y_pred = cnn.predict(X_test)
        y_pred[5]

313/313 — 1s 3ms/step

Out[26]:array([[3.13081604e-04, 4.44748948e-05, 5.50113909e-04, 7.99367249e-01,
8.86134439e-05, 4.03291918e-02, 1.07765093e-01, 6.09774315e-06,
5.11943251e-02, 3.41743929e-04],
[5.92878573e-02, 1.65216730e-03, 1.35267464e-05, 9.33361566e-09,
8.25674178e-08, 2.40147768e-09, 1.99592201e-08, 1.20279198e-09,
9.38898563e-01, 1.47697487e-04],
[5.96161075e-02, 1.73812702e-01, 2.15590112e-02, 7.67891994e-03,
1.13394135e-03, 3.51405656e-03, 2.92896270e-03, 2.28178990e-03,
7.09434368e-01, 1.80395674e-02],
[9.83872950e-01, 4.68066632e-04, 6.39343983e-04, 1.89955153e-05,
4.20502511e-05, 7.97551638e-06, 4.19331955e-05, 9.58110468e-06,
1.47862770e-02, 1.12778129e-04],
[7.27530676e-07, 4.49761183e-06, 7.50960410e-03, 1.68715734e-02,
4.27279890e-01, 6.18672639e-04, 5.47628284e-01, 1.12274364e-07,
8.65199254e-05, 7.82810829e-08]], dtype=float32)
```


```
In [27]:y_classes = [np.argmax(element) for element in y_pred]
        y_classes[5]

Out[27]:[3, 8, 8, 0, 6]

In [28]:y_test[5]

Out[28]:array([3, 8, 8, 0, 6], dtype=uint8)

In [29]:plot_sample(X_test, y_test, 60)
```




```
In [30]:plot_sample(X_test, y_test, 100)

In [31]:plot_sample(X_test, y_test, 100)

In [32]:plot_sample(X_test, y_test, 100)

In [33]:plot_sample(X_test, y_test, 79)
```



```
In [34]:classes[y_classes[79]]

Out[34]: 'ship'

In [35]:

Out[35]:
```